

Two Assistants That Know What The Other Heard

HR gets a second chatbot for project requirements documents, separate from the people-search one. Each can see what was discussed in the other, and use it to suggest the next useful step. **Version 3** — corrects a wrong reading of the design in earlier drafts.

The problem

What a project needs — the skills, the constraints, the caveats — lives in a Word or PDF file nobody reads twice. Meanwhile HR is searching the staff directory to find people for that same project. The two activities never meet, so the same questions get asked twice and the answers don't carry over.

What HR gets

- **A PRDs page** — upload a project document, see the requirements it found as an editable list, fix anything wrong, save.
- **Its own chatbot** on that page, which only answers about requirements. It cannot search for people.
- **The existing search chatbot**, unchanged, for finding people.
- **Each one aware of the other.** Ask about staffing after capturing Meridian's requirements, and the search assistant offers to check who covers them. Ask about Meridian's needs after a fruitless Terraform search, and the PRD assistant points out the overlap.
- **Conversations that survive.** Close the tab, log back in tomorrow, both threads are still there.

The three rules that keep this safe

Only facts cross between them	Not the transcript. When one assistant sees the other's history, it gets a short list like "project discussed: Meridian; skill discussed: Terraform" — never the sentences themselves, and never anything the AI wrote. Those are project and skill names from our own controlled list, so uploaded document text has no route into the other chatbot.
Only HR and the project owner read the notes	An earlier draft locked the page and the chat but left the plain "list this project's notes" address open to any employee — a lock on the front door with a window open. Notes are sentences copied out of a planning document, so they now use the same lock. Plain skill lists stay broadly readable: those are ordinary company facts.
Note text never reaches the AI when it writes an answer	Reading an uploaded file needs the whole document — that is the job, and HR just chose the file. Writing the one-line summary afterwards only needs the skill names, so notes are left out. The reason is less about privacy than about trust: an uploaded document can be written by anyone, including someone outside the company, and a sentence in it could be an instruction aimed at the AI. HR still sees notes on screen.

Build order

1 Storage	A table for the notes and the routes to read and write them, with the new lock. No AI, no screens.
2 Reading documents	Pull requirements out of an uploaded file; erase the file's text once HR confirms.
3 Remembering	Save both conversations so they survive a reload. Nothing cross-surface yet.
4 The second chatbot	Its own address, its own small tool set, its own limits. The search one is untouched.
5 The screen	The "PRDs" tab: pick project, upload, edit, confirm, chat.
6 The link	Each assistant reads facts from the other and offers a follow-up. Built last, on proven parts.

Worth knowing before starting

This is bigger than the last draft. Saving conversations was cut early on to keep scope down, and wanting them to survive a re-login brings it back — that is phase 3, and it is real work. It also removes the fiddliest part of the previous plan: because the PRD assistant now has its own address, there is no need to filter one shared tool list by who is asking. Two other things fall out nicely — saved conversations reuse a storage shape the code already has and already treats carefully, and the second assistant becomes the first real test of the per-assistant limits built earlier this week, which until now had exactly one entry in them.